

Deep Q-Network (DQN)



Was kann DQN?



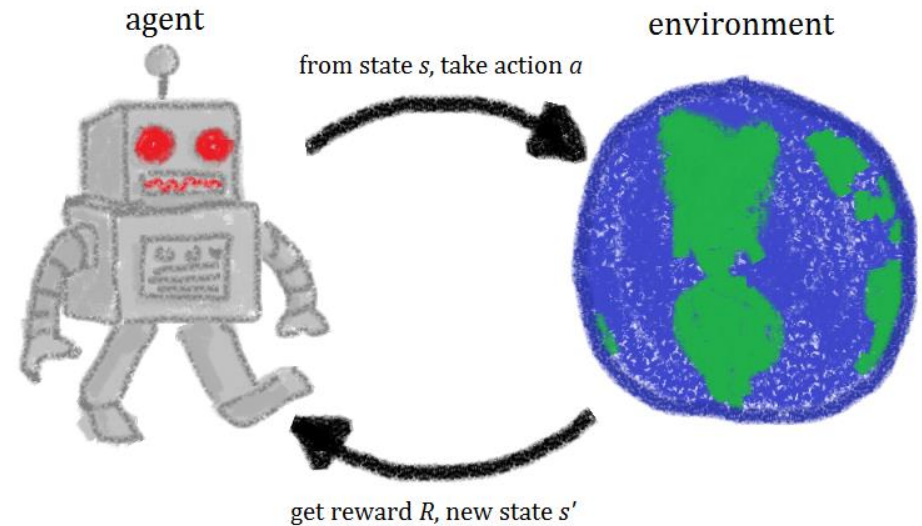
Breakout



Was ist DQN?

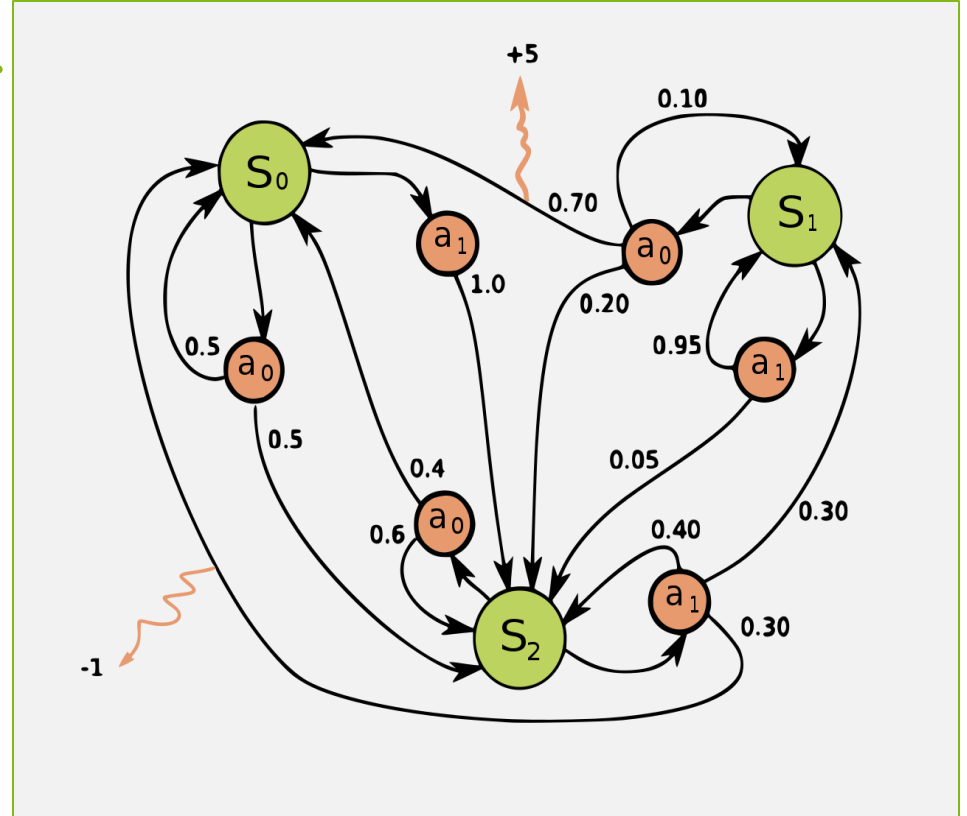
Agent-Environment Loop

- Agent **perceives** the **state** of the environment
- Agent **takes** an **action**
- Environment **transitions** to a **new state**
- Environment **provides** a **reward** to the agent



Markov Decision Process..

- *States*(S)
- *Actions*(A)
- *transition Function*:
 $P(s'|s, a)$
- *Reward Function*:
 $R(s, a, s')$



.. and its return

The return G_t is the total accumulated reward from timestep t onwards.

Defined as:

$$G_t = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

Policy

A Policy π is a function that defines the **behavior** of an agent by specifying the **probability** of taking each action a in state s

$$\pi(a|s) = P(a_t = a | s_t = s)$$

The goal of a policy is to **maximize** the expected return starting from each state

Action value function

Action value function $q_{\pi}(s, a)$ under policy π is the **expected return** after taking action a in state s following π

$$q_{\pi}(s, a) = E_{\pi}[G_t | s_t = s, a_t = a]$$

Bellman's expectation equation..

.. for action value functions:

$$q_{\pi}(s, a) = E[R_{t+1} + \gamma q_{\pi}(S_{t+1}, A_{t+1}) | s_t = s, a_t = a]$$

Bellman's optimality equation..

.. for action value functions:

$$q_*(s, a) = E[R_{t+1} + \gamma \max_{a'} q_*(S_{t+1}, a') | s_t = s, a_t = a]$$

Where q_* is the optimal action value function.

Q-Learning

- Model-free reinforcement learning algorithm
- Aims to find the optimal action value function q_*
- Update rule is:
$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha [R_{t+1} + \gamma \max_a Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t)]$$
$$Q(s_t, a_t) \leftarrow (1 - \alpha) * Q(s_t, a_t) + \alpha [R_{t+1} + \gamma \max_a Q(s_{t+1}, a_{t+1})]$$
- Related to bellman optimality equation
- Iteratively updates action value function to converge to q_*

Epsilon-Greedy Policy

- Balances **exploration** and **exploitation** by choosing a **random** action
- Random action with ϵ and best-known action with $1 - \epsilon$
- Helps ensure **sufficient exploration** of the state-action space
- Crucial for convergence of Q-Learning to optimal action value function

Einführung in das Deep Q Learning

Policy Gradients



Go Right

??

Deep Q-Learning



Please wait, I am still
calculating Q value, only
41891 actions left...

update rule & loss function

- Update rule

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha [R_{t+1} + \gamma \max_a Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t)]$$

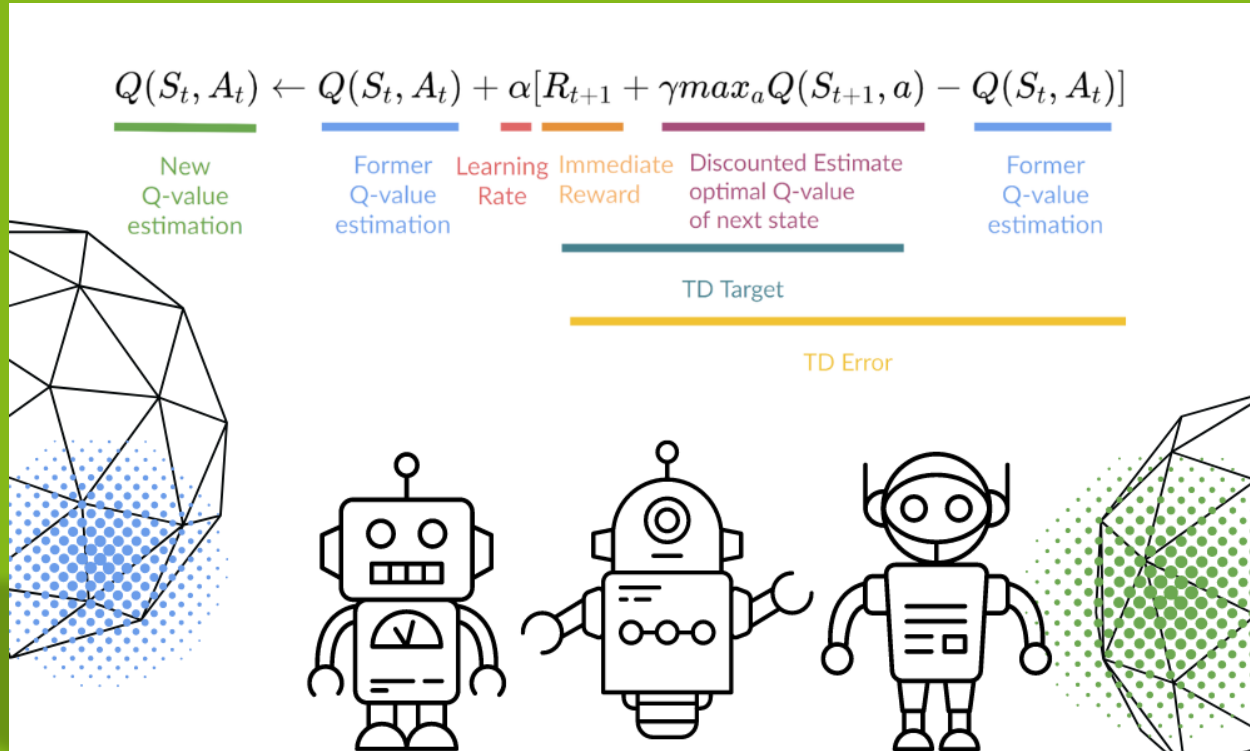
- Loss function

$$L_i(\theta_i) = \mathbb{E}_{s, a \sim p(\cdot)} [(y_i - Q(s_t, a_t; \theta_i))^2]$$

$$y_i = [r + \gamma \max_{a_{t+1}} Q(s_{t+1}, a_{t+1}; \theta_{i-1})]$$

$$L_i(\theta_i) = \mathbb{E}_{s, a \sim p(\cdot)} [((r + \gamma \max_{a_{t+1}} Q(s_{t+1}, a_{t+1}; \theta_{i-1})) - Q(s_t, a_t; \theta_i))^2]$$

Update rule a closer look!



Gradient

$$\nabla_{\theta_i} L_i(\theta_i) = \mathbb{E}_{s,a \sim p(\cdot)} [(y_i - Q(s_t, a_t; \theta_i)) \nabla_{\theta_i} Q(s_t, a_t; \theta_i)]$$

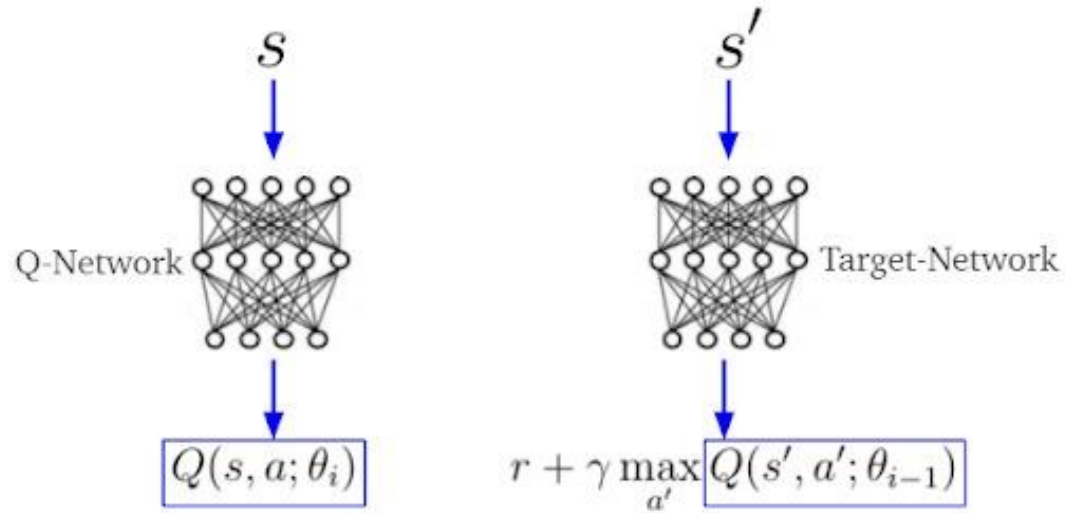
$$\nabla_{\theta_i} L_i(\theta_i) = \mathbb{E}_{s,a \sim p(\cdot)} \left[\left(r + \gamma \max_{a_{t+1}} Q(s_{t+1}, a_{t+1}; \theta_{i-1}) - Q(s_t, a_t; \theta_i) \right) \nabla_{\theta_i} Q(s_t, a_t; \theta_i) \right]$$

- to minimize the loss function
- computing the gradient for updates: $\theta_{i+1} \leftarrow \theta_i - \alpha \nabla_{\theta_i} L_i(\theta_i)$

Experience Replay

- Memory of experiences
- Pick of minibatches for update of the Q-function
- Addresses three issues:
 - Data Efficiency
 - Low Correlation
 - Catastrophic forgetting

Target Q-Network



DQN Pseudocode

Algorithm 1 Deep Q-learning with Experience Replay

Initialize replay memory \mathcal{D} to capacity N

Initialize action-value function Q with random weights

for episode = 1, M **do**

 Initialize sequence $s_1 = \{x_1\}$ and preprocessed sequenced $\phi_1 = \phi(s_1)$

for $t = 1, T$ **do**

 With probability ϵ select a random action a_t

 otherwise select $a_t = \max_a Q^*(\phi(s_t), a; \theta)$

 Execute action a_t in emulator and observe reward r_t and image x_{t+1}

 Set $s_{t+1} = s_t, a_t, x_{t+1}$ and preprocess $\phi_{t+1} = \phi(s_{t+1})$

 Store transition $(\phi_t, a_t, r_t, \phi_{t+1})$ in \mathcal{D}

 Sample random minibatch of transitions $(\phi_j, a_j, r_j, \phi_{j+1})$ from \mathcal{D}

 Set $y_j = \begin{cases} r_j & \text{for terminal } \phi_{j+1} \\ r_j + \gamma \max_{a'} Q(\phi_{j+1}, a'; \theta) & \text{for non-terminal } \phi_{j+1} \end{cases}$

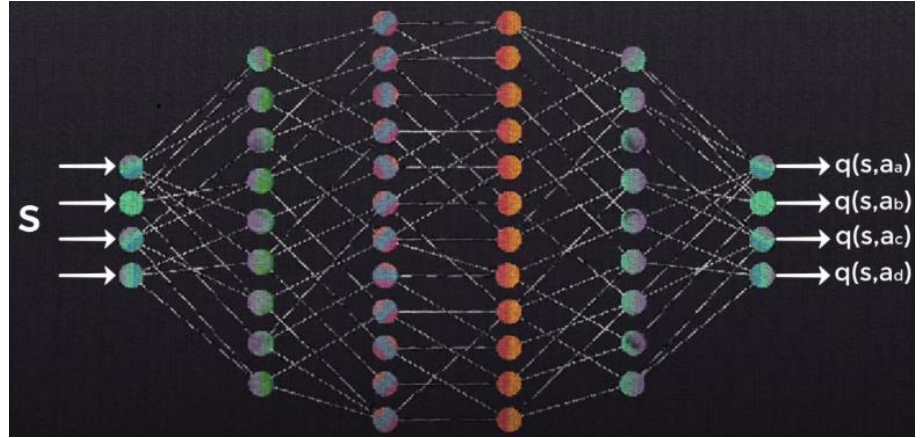
 Perform a gradient descent step on $(y_j - Q(\phi_j, a_j; \theta))^2$ according to equation 3

end for

end for

Neuronal Network

- Q-function » Q-network » neuronal network
- Weights of q-learning are weights for neuronal network
- Layers: input (states), hidden (convolutional layers), output (q-values for each action)



<https://youtu.be/xVkPh9E9GfE>

Preprocessing of game frames

1. Grayscale
2. Downsampling to 110x84
3. Crop to 84x84
4. Stack 4 frames to one

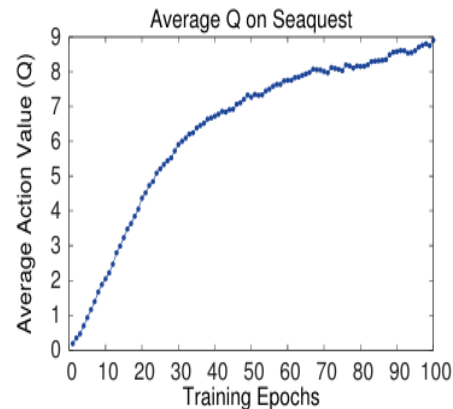
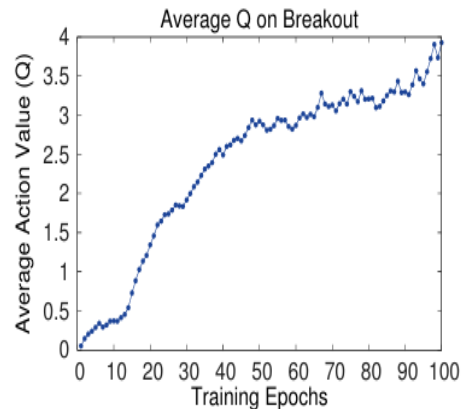
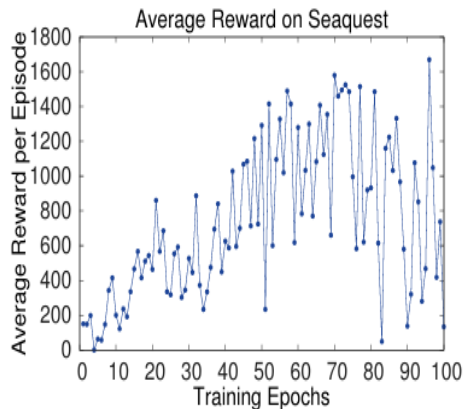
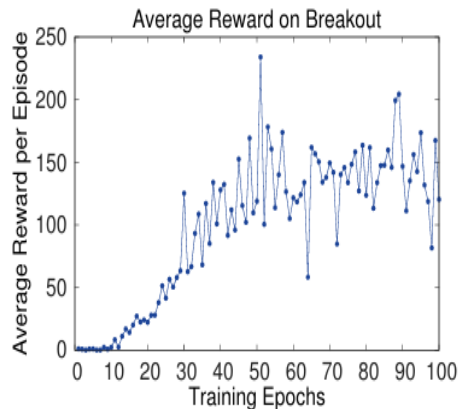
Reward Clipping

- limit the reward to interval $[-1,1]$
- High rewards are counterproductive for strabelized learning
- Clipping stabelizes learning

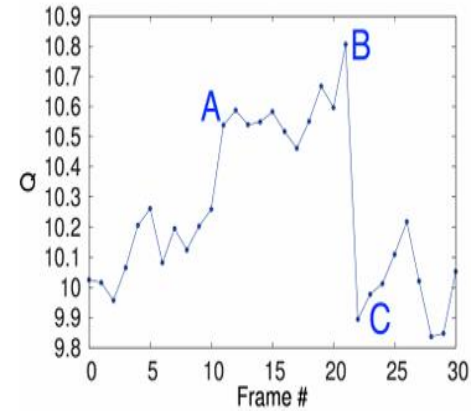
Performance Metrics

- Within a specific time periode
 - Average and maximum points

Results from the Paper



Results from the Paper



How to implement

Classes:

Agent Class (DQNAgent):

- Responsible for controlling the DQN agent.
- Methods:
 - `select_action`: Selects actions based on the current state.
 - `train`: Trains the agent using experiences stored in the replay memory.
 - `update_target_network`: Updates the target network periodically.

How to implement

Classes:

Neural Network Class (DQN):

- Defines the neural network architecture for approximating Q-values.
- Methods:
 - `__init__`: Initializes the network with `state_size` and `action_size`.
 - `forward`: Performs forward pass computation.
 - `compute_loss`: Computes the loss between predicted and target Q-values.
 - `backward`: Performs gradient descent.

How to implement

Classes:

Replay Memory Class:

- Manages the replay memory buffer.
- Methods:
 - `add_experience`: Adds experiences to the replay memory.
 - `sample_batch`: Samples batches of experiences for training.

Links / Sources

Picture Roboter & World :

https://commons.wikimedia.org/wiki/File:RL_agent.png

Markov decision process:

https://upload.wikimedia.org/wikipedia/commons/thumb/a/ad/Markov_Decision_Process.svg/1280px-Markov_Decision_Process.svg.png

Update rule:

https://images.datacamp.com/image/upload/v1666973295/Q_learning_equation_3cd6652b98.png

Policy bears:

<https://pylessons.com/media/Tutorials/Reinforcement-learning-tutorial/Beyond-DQN/Beyond-DQN.jpg>

Target Q-network :

https://builtin.com/sites/www.builtin.com/files/styles/ckeditor_optimize/public/inline-images/4_double-deep-q-learning.png

Results from the paper:

[Volodymyr Mnih, Koray Kavukcuoglu, David Silver, Alex Graves, Ioannis Antonoglou, Daan Wierstra, Martin Riedmiller, Playing atari with deep reinforcement learning \(2013\),
https://arxiv.org/pdf/1312.5602](https://arxiv.org/pdf/1312.5602)